



Multiclass classification:

- (a) Download the CIFAR-10² and load the pickled data into your program.
- (b) Build a basic convolutional neural networks with several convolution, pooling, and normalization layers. Flatten the output of the convolution layers and pass it to a single dense layer that will produce the output using softmax activation.
- (c) Test the performance of the model you built and tune hyper parameters as needed.
- (d) Add one or two inception blocks and test performance.
- (e) Remove the inception blocks and add one or two residual blocks instead. Test performance and compare to previous results.

1. Hyperparameter tuning:

- i. Initial model with 2 convolution layers, max pooling layers, Adding dropout rate 0.2:
Model summary:

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	(None, 32, 32, 3)	0
conv2d_7 (Conv2D)	(None, 32, 32, 64)	256
leaky_re_lu_6 (LeakyReLU)	(None, 32, 32, 64)	0
max_pooling2d_6 (MaxPooling2	(None, 16, 16, 64)	0
dropout_7 (Dropout)	(None, 16, 16, 64)	0
batch_normalization_6 (Batch	(None, 16, 16, 64)	256
conv2d_8 (Conv2D)	(None, 16, 16, 128)	73856
leaky_re_lu_7 (LeakyReLU)	(None, 16, 16, 128)	0
max_pooling2d_7 (MaxPooling2	(None, 8, 8, 128)	0
dropout_8 (Dropout)	(None, 8, 8, 128)	0
batch_normalization_7 (Batch	(None, 8, 8, 128)	512
flatten_2 (Flatten)	(None, 8192)	0
dropout_9 (Dropout)	(None, 8192)	0
dense_2 (Dense)	(None, 10)	81930
Total params: 156,810		
Trainable params: 156,426		
Non-trainable params: 384		



2. Adding inception blocks:

```
inception1 = Conv2D(32, kernel_size=1, strides=1, input_shape=self.img_shape, padding
= 'same')(img)
```

```
hid1 = LeakyReLU()(inception1)
hid1 = MaxPooling2D((2,2))(hid1)
hid1 = Dropout(0.2)(hid1)
hid1 = BatchNormalization()(hid1)
print(hid1)
```

```
inception2 = Conv2D(32, kernel_size=3, strides=1, input_shape=self.img_shape,
padding = 'same')(img)
```

```
hid2 = LeakyReLU()(inception2)
hid2 = MaxPooling2D((2,2))(hid2)
hid2 = Dropout(0.2)(hid2)
hid2 = BatchNormalization()(hid2)
print(hid2)
```

```
inception3 = Conv2D(32, kernel_size=5, strides=1, input_shape=self.img_shape,
padding = 'same')(img)
```

```
hid3 = LeakyReLU()(inception3)
hid3 = MaxPooling2D((2,2))(hid3)
hid3 = Dropout(0.2)(hid3)
hid3 = BatchNormalization()(hid3)
print(hid3)
```

Model: "model_3"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_4 (InputLayer)	(None, 32, 32, 3)	0	

conv2d_9 (Conv2D)	(None, 32, 32, 32)	128	input_4[0][0]

conv2d_10 (Conv2D)	(None, 32, 32, 32)	896	input_4[0][0]

conv2d_11 (Conv2D)	(None, 32, 32, 32)	2432	input_4[0][0]

leaky_re_lu_8 (LeakyReLU)	(None, 32, 32, 32)	0	conv2d_9[0][0]

leaky_re_lu_9 (LeakyReLU)	(None, 32, 32, 32)	0	conv2d_10[0][0]

leaky_re_lu_10 (LeakyReLU)	(None, 32, 32, 32)	0	conv2d_11[0][0]

max_pooling2d_8 (MaxPooling2D)	(None, 16, 16, 32)	0	leaky_re_lu_8[0][0]

max_pooling2d_9 (MaxPooling2D)	(None, 16, 16, 32)	0	leaky_re_lu_9[0][0]

max_pooling2d_10 (MaxPooling2D)	(None, 16, 16, 32)	0	leaky_re_lu_10[0][0]

dropout_10 (Dropout)	(None, 16, 16, 32)	0	max_pooling2d_8[0][0]


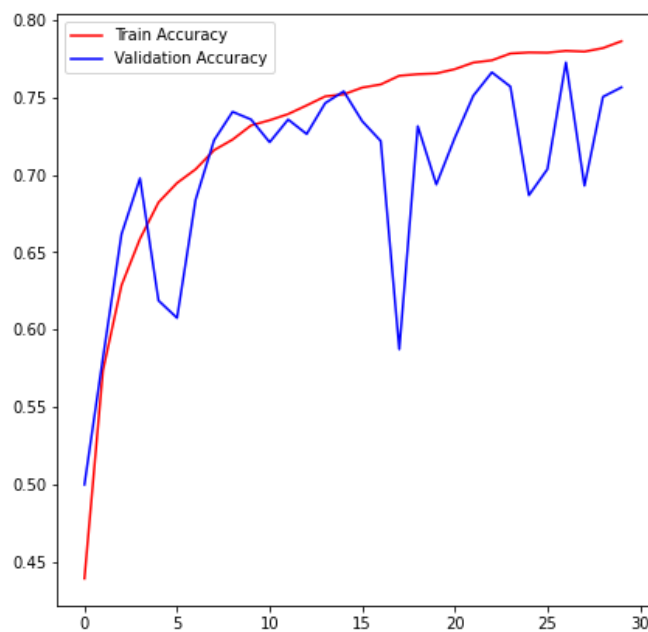
dropout_11 (Dropout)	(None, 16, 16, 32)	0	max_pooling2d_9[0][0]

dropout_12 (Dropout)	(None, 16, 16, 32)	0	max_pooling2d_10[0][0]

batch_normalization_8	(BatchNormalizer)	(None, 16, 16, 32)	128	dropout_10[0]
batch_normalization_9	(BatchNormalizer)	(None, 16, 16, 32)	128	dropout_11[0]
batch_normalization_10	(BatchNormalizer)	(None, 16, 16, 32)	128	dropout_12[0]
concatenate_2	(Concatenate)	(None, 16, 16, 96)	0	batch_normalization_8[0][0]
				batch_normalization_9[0][0]
				batch_normalization_10[0][0]
conv2d_12	(Conv2D)	(None, 16, 16, 64)	6208	concatenate_2[0][0]
leaky_re_lu_11	(LeakyReLU)	(None, 16, 16, 64)	0	conv2d_12[0]
max_pooling2d_11	(MaxPooling2D)	(None, 8, 8, 64)	0	leaky_re_lu_11[0][0]
dropout_13	(Dropout)	(None, 8, 8, 64)	0	max_pooling2d_11[0][0]
batch_normalization_11	(BatchNormalizer)	(None, 8, 8, 64)	256	dropout_13[0]
conv2d_13	(Conv2D)	(None, 8, 8, 128)	73856	batch_normalization_11[0][0]
leaky_re_lu_12	(LeakyReLU)	(None, 8, 8, 128)	0	conv2d_13[0]
max_pooling2d_12	(MaxPooling2D)	(None, 4, 4, 128)	0	leaky_re_lu_12[0][0]
dropout_14	(Dropout)	(None, 4, 4, 128)	0	max_pooling2d_12[0][0]

batch_normalization_12 (Batch Normalization)	(None, 4, 4, 128)	512	dropout_14[0][0]
flatten_3 (Flatten)	(None, 2048)	0	batch_normalization_12[0][0]
dropout_15 (Dropout)	(None, 2048)	0	flatten_3[0]
dense_3 (Dense)	(None, 10)	20490	dropout_15[0][0]

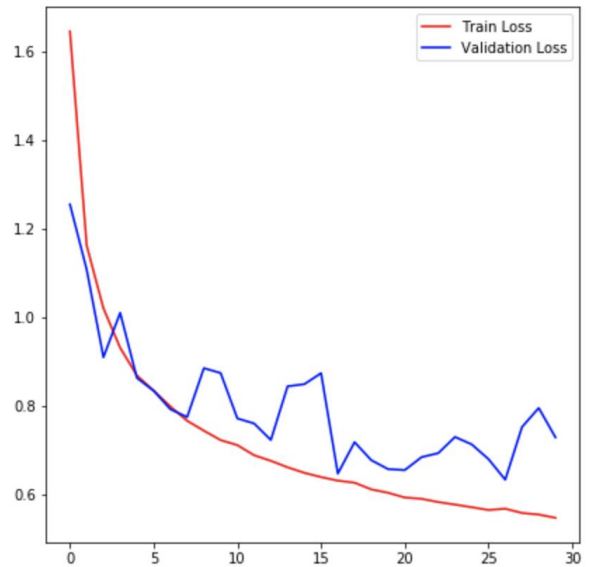

Total params: 105,162
 Trainable params: 104,586
 Non-trainable params: 576



Adding inception blocks has increased the performance to 75% on the validation set. This is because inception blocks increase the dimensionality of the features extracted by the convolution blocks.

3. Residual Block

- Adding residual blocks has a better performance than the inception blocks.



Adding a second residual block

- A second residual block has not improved the performance, and is not necessary for our use case.

